

Getting Started with Kubernetes on **AWS**

Brought to you by the AWS Sydney Cloud Support Team

Day 2

Agenda

- Kubernetes basic resources
- Exploring the cluster with kubectl
- Designing your application for high availability
- Architecting your application using services
- Exposing your application to the world

Firstly...

Do you have a cluster?

Using a terminal in Cloud9, verify there are Worker Nodes in your cluster

```
$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
ip-192-168-29-29.ap-southeast-1.compute.internal	Ready	<none>	7d	v1.13.7
ip-192-168-88-24.ap-southeast-1.compute.internal	Ready	<none>	7d	v1.13.7

[Link to GIF](#)

Do you have a cluster?

Using a terminal in Cloud9, verify there are Worker Nodes in your cluster

```
$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
ip-192-168-29-29.ap-southeast-1.compute.internal	Ready	<none>	7d	v1.13.7
ip-192-168-88-24.ap-southeast-1.compute.internal	Ready	<none>	7d	v1.13.7

[Link to GIF](#)

Ooops, no, I don't have a cluster!

<https://github.com/aws-els/eks>

```
eksctl create cluster --ssh-access --version 1.13 --node-type t3.medium --name eks
```

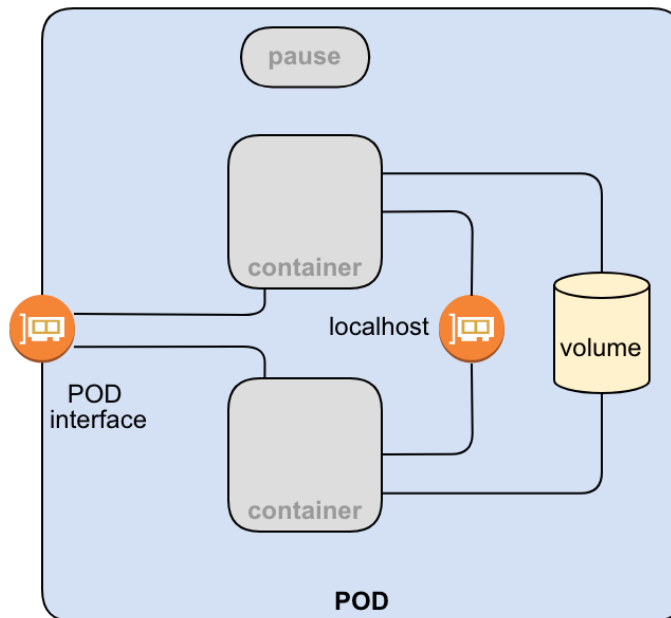
Lets pull in the latest changes

```
$ cd eks  
$ git pull
```

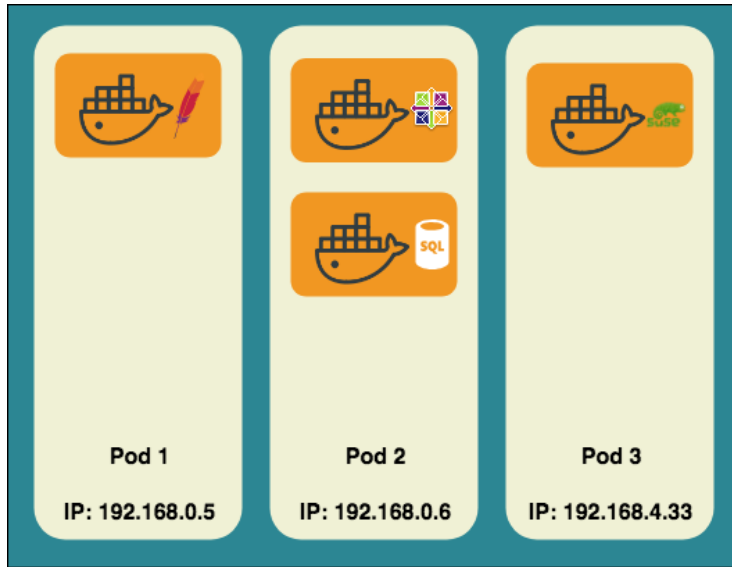
Pod

What is a Pod?

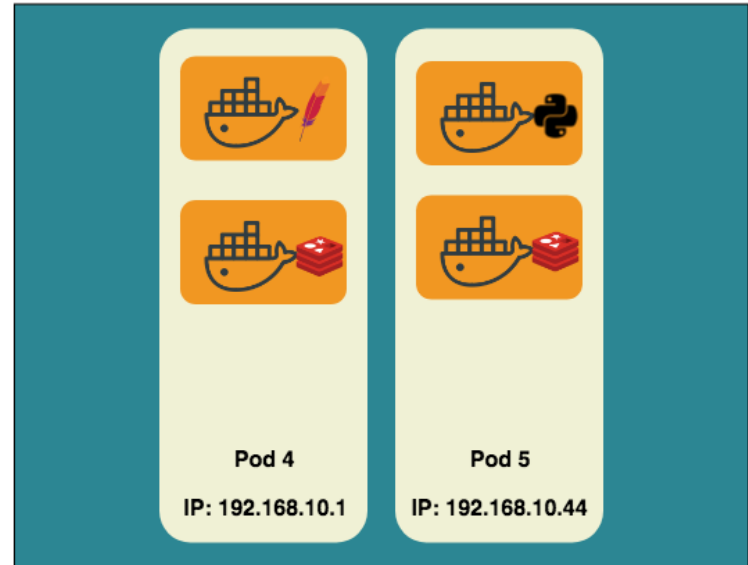
- The smallest building block of Kubernetes
- A Pod encapsulates the container(s) and resources needed to run the application
- A unit of deployment



What is a Pod?



Worker Node 01



Worker Node 02

Creating Pods in Kubernetes

Lab 1: Define a Pod

Pod definition

```
apiVersion: v1
kind: Pod
metadata:
  name: web-server
spec:
  containers:
  - name: container1
    image: nginx
```

Lab 1: Define a Pod

Pod definition

```
apiVersion: v1
kind: Pod
metadata:
  name: web-server
spec:
  containers:
  - name: container1
    image: nginx
```

File: labs/01_pod.yaml

Lab 1: Creating a Pod

Send the definition to the cluster:

```
kubectl create -f labs/01_pod.yaml
```



Lab 1: Check the Pod

List and describe the Pod

```
kubectl get pod
```


```
kubectl describe pod web-server
```

Lab 1: Check the Pod

Create a tunnel and connect to your Pod

```
kubectl port-forward pod/web-server 8080:80 &  
curl localhost:8080
```

Or, connect using a browser and Cloud9



Lab 1: Clean up

We ran the port-forward in the background, lets clean it up before we move on

In the Cloud9 terminal to bring it back to the foreground:

```
fg
```

```
# Hit Ctrl + C to kill the port-forward
```

Lab 2: Working with Pods

Check the logs

```
kubectl logs web-server
```

Lab 2: Working with Pods

Check the logs

```
kubectl logs web-server
```

Connect

```
# run one command  
kubectl exec web-server cat /etc/hostname  
  
# run with console connected to pod  
kubectl exec -it web-server -- bash
```

Lab 2: Working with Pods

Check the logs

```
kubectl logs web-server
```

Connect

```
# run one command  
kubectl exec web-server cat /etc/hostname  
  
# run with console connected to pod  
kubectl exec -it web-server -- bash
```

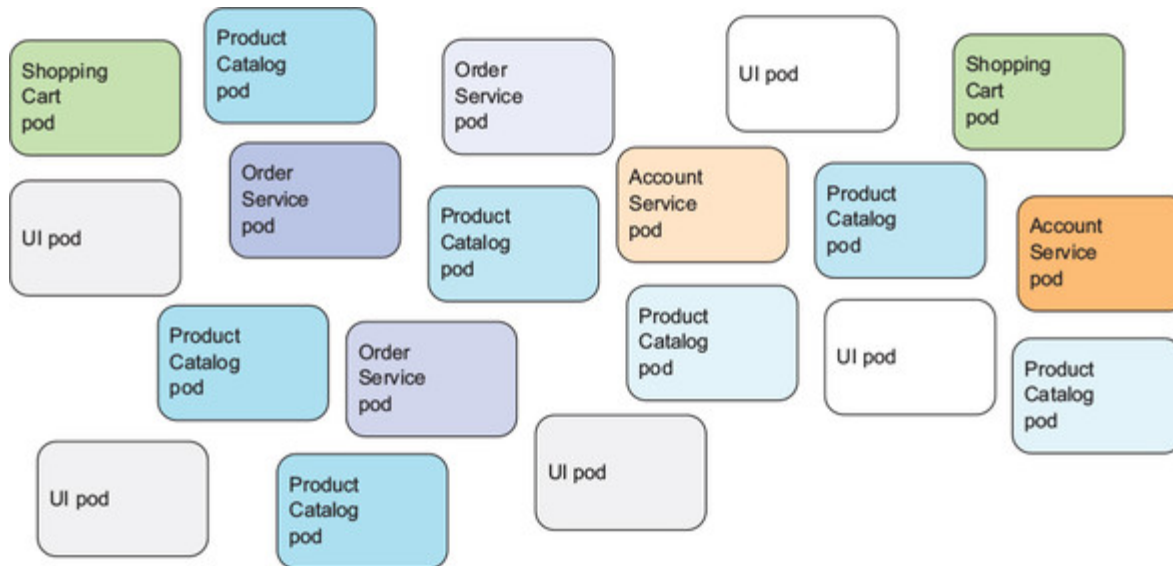
You can delete it: *But keep it for now!*

```
kubectl delete pod web-server  
  
# OR  
  
kubectl delete -f labs/01_pod.yaml
```

Labels

Labels

What if we are running a lot of pods?

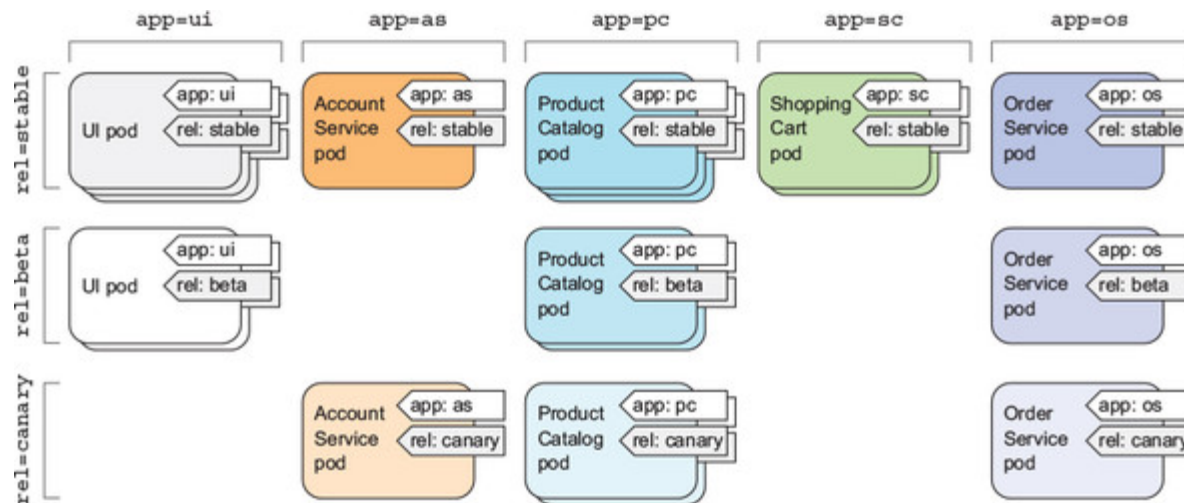


Picture from [Kubernetes in action](#)

Labels

Labels are key/value pair tags (like tags in AWS)

Can be used to query and organize resources



Picture from [Kubernetes in action](#)

Lab 3: Applying labels

Labels in the object definition

```
apiVersion: v1
kind: Pod
metadata:
  name: echo-server
  labels:
    env: training
    type: single_pod
spec:
  containers:
  - name: echo
    image: k8s.gcr.io/echoserver:1.4
```

Checking the labels

```
kubectl apply -f labs/03_labels.yaml

kubectl get pods --show-labels
```

File: labs/03_labels.yaml

Links

Pods

- <https://kubernetes.io/docs/concepts/workloads/pods/pod-overview/>
- <https://kubernetes.io/docs/concepts/workloads/pods/pod/>

Kubernetes in Action

- <https://www.safaribooksonline.com/library/view/kubernetes-in-action/9781617293726/>

Demos

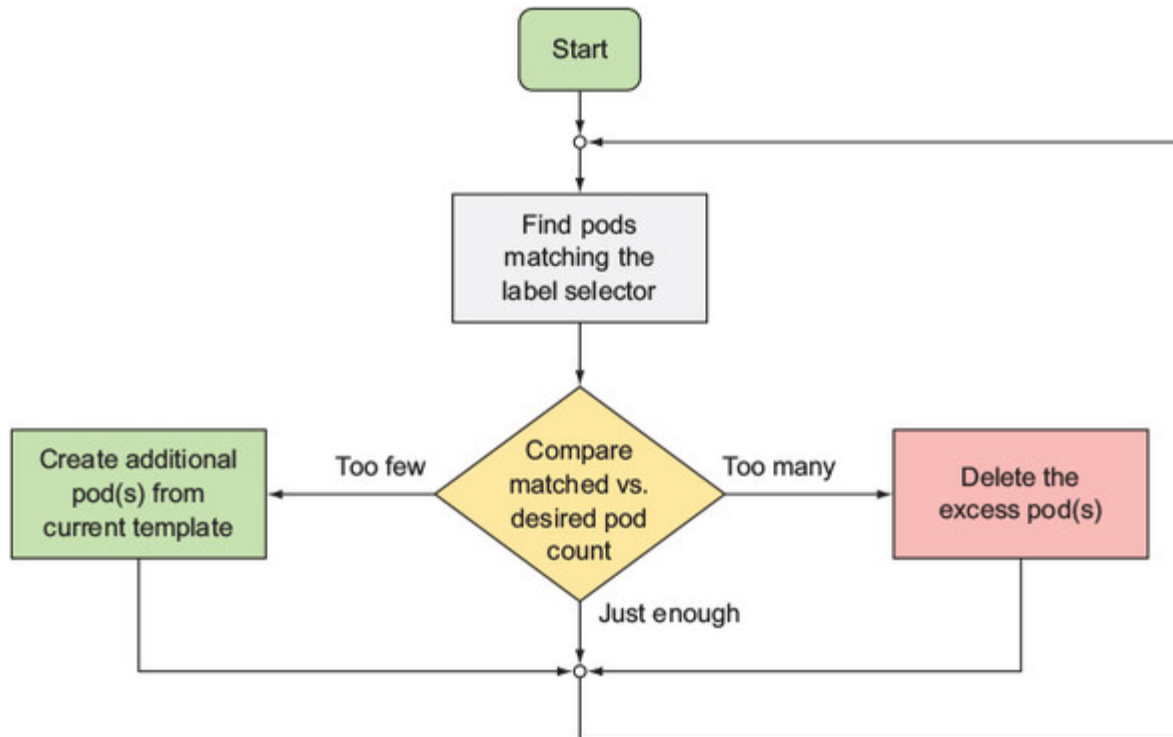
- <https://eksworkshop.com/>
- <https://github.com/kubernetes/contrib/blob/master/micro-demos/>

Controllers

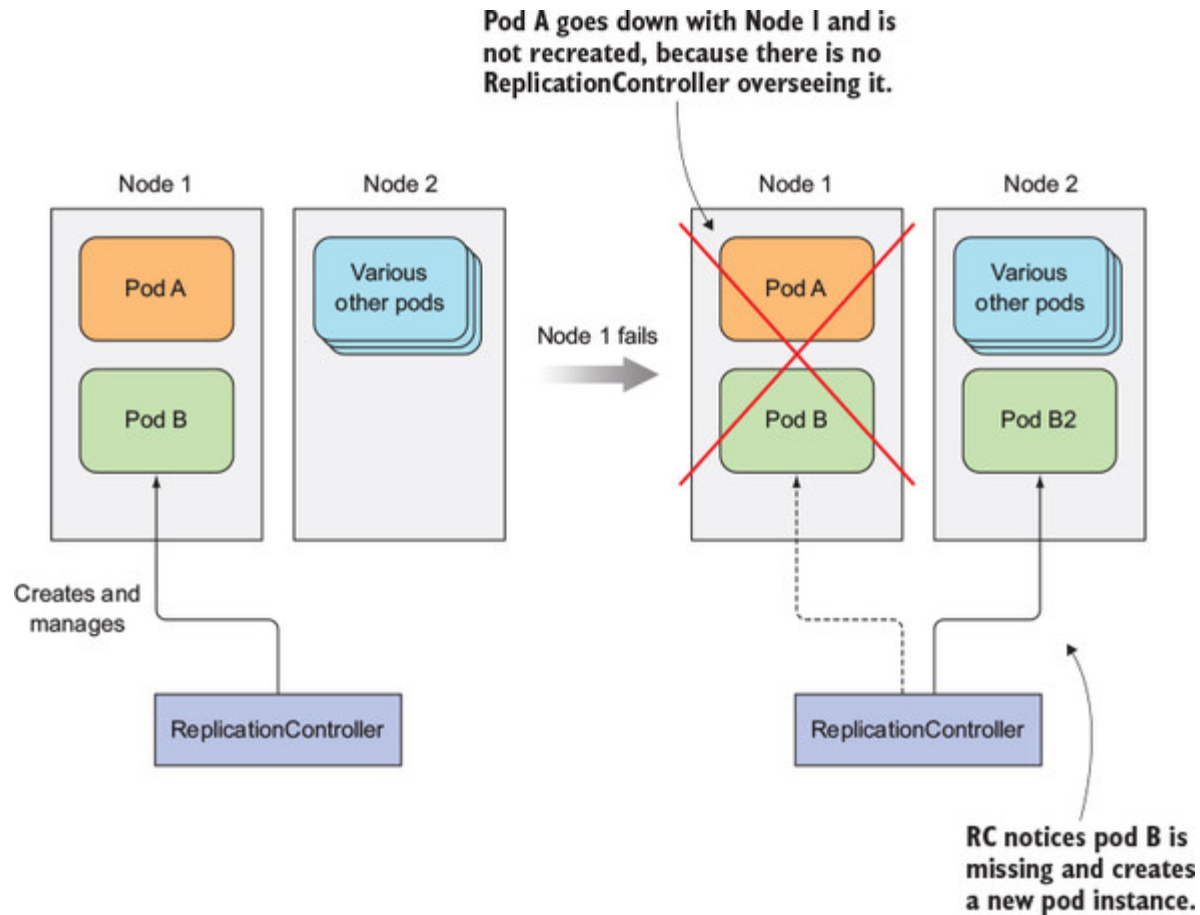
What controllers are and what they do?

- Resource responsible for managing pods
- Ensure pods are always running
- Replace missing and unhealthy pods
- Delete 'extra' pods
- Provide an easy way to scale the application
- Rely on Labels to account for the pods

What controllers are and what they do?



What controllers are and what they do?



Controllers

- Most commonly used controllers in Kubernetes:
 - **ReplicationController**
 - **ReplicaSet** the next generation of Replication Controllers
 - **Deployments** - preferred way to manage Replica Sets
 - **DaemonSet**
 - **Jobs**
 - **CronJobs**
 - **StatefulSets**

Deployments

Deployments

What is it?

- A Deployment controller provides declarative updates at controlled rate for Pods
- An easy way to deploy updates for existing applications
- Allows you to pause/resume deployments

Deployments

Comparing: Pods vs Deployment:

Single Pod:

```
apiVersion: v1
kind: Pod
metadata:
  name: web-server
  labels:
    app: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.7.9
```

Deployment:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: web-server-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
```

File: labs/04_deployment.yaml

Deployments

Lab 4: Create a Deployment

- Creating a Deployment

```
kubectl apply -f labs/04_deployment.yaml
```

- Checking the results

```
kubectl get deployments  
kubectl get pod  
kubectl get pod -l app=nginx
```

- Scale out the Deployment

```
kubectl scale deployment web-server-deployment --replicas=5
```

- Check the nginx server version

```
kubectl port-forward web-server-deployment-XXXXXX-YYYYY 8080:80
```

Deployments

Lab 4: Scale the Deployment

- Scale in the deployment. Let's be frugal

```
kubectl edit deployment web-server-deployment  
  
# set spec.replicas to 3
```

- Checking the results

```
kubectl get pod  
  
# OR  
  
kubectl get pod -l app=nginx -L app
```

Deployments

Lab 4: Update the Deployment

- Update the deployment

```
kubectl set image deployment web-server-deployment nginx=httpd
```

```
# OR
```

```
kubectl edit deployment web-server-deployment  
# change spec.template.spec.containers.image to httpd
```

- Checking the results

```
kubectl rollout status deployment web-server-deployment
```

```
# OR
```

```
kubectl get pod
```

```
# OR
```

```
kubectl get pod -l app=nginx -L app
```

- Check the Nginx server version now

```
kubectl port-forward web-server-deployment-XXXXXX-YYYYY 8080:80
```

Services

Services

What's a service?

- Service is another layer on top of the pods
- Instead of connecting to the pods directly we connect to the service instead
- Very similar to a load balancer

Services

What's a service?

- Service is another layer on top of the pods
- Instead of connecting to the pods directly we connect to the service instead
- Very similar to a load balancer

But why?

- Pods are ephemeral
- Pods' IPs are dynamic
- A single application might contain several Pods

Services

What's a service?

- Service is another layer on top of the pods
- Instead of connecting to the pods directly we connect to the service instead
- Very similar to a load balancer

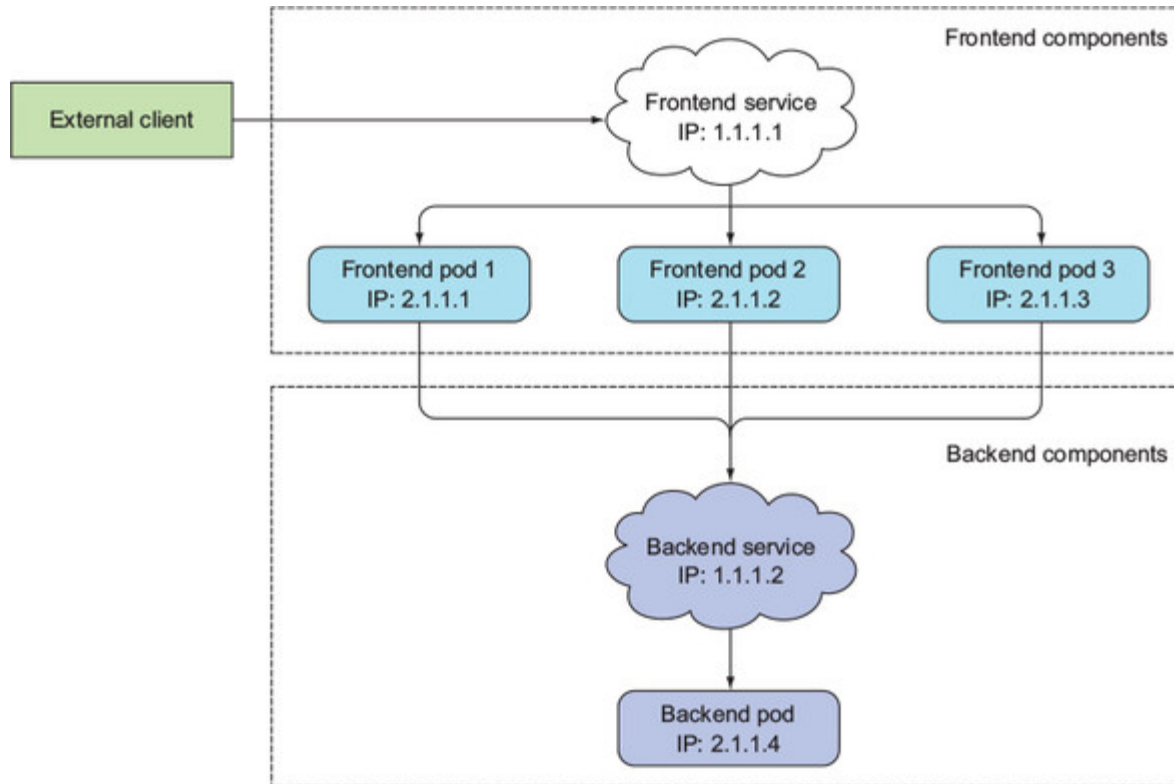
But why?

- Pods are ephemeral
- Pods' IPs are dynamic
- A single application might contain several Pods

So, how to reach an application?

Services

A **Service** is an abstraction layer which enables external traffic exposure, load balancing and service discovery



* from <https://kubernetes.io/>

Lab 5: Create a Service

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-app
spec:
  ports:
  - port: 80
    targetPort: 80
  selector:
    app: nginx
```

File: labs/05_services.yaml

Lab 5: Create a Service

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-app
spec:
  ports:
  - port: 80
    targetPort: 80
  selector:
    app: nginx
```

- Creating a service

```
kubectl apply -f labs/05_services.yaml
```

File: labs/05_services.yaml

Lab 5: Create a Service

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-app
spec:
  ports:
    - port: 80
      targetPort: 80
  selector:
    app: nginx
```

- Creating a service

```
kubectl apply -f labs/05_services.yaml
```

- Checking the results

```
kubectl get services
kubectl describe svc nginx-app
```

File: labs/05_services.yaml

Services

- There are different types of services:
 - **ClusterIP** is the default. Used for intra-cluster communication
 - **LoadBalancer** provisions a Load Balancer for you.

Connecting to your service:

```
kubectl edit svc nginx-app  
  
# change spec.type from 'ClusterIP' to 'LoadBalancer'
```

- Checking the results

```
kubectl get svc nginx-app
```

- Now you should get the LB URL from the EC2 console and open it in your browser.

End of Day 2

Questions?

Thank you!